

# **Blueprint for Deterministic Edge Services: Cloudflare Workers, KV, and Scheduled Events**

Author: Brian Barry, SETU, C00274624@gmail.com

## Contents

Abstract .....	3
1 Introduction.....	4
2 Related Work.....	5
3 System Architecture .....	6
4 Implementation Details.....	8
5 Performance Evaluation.....	10
6 Limitations and Future Work .....	11
7 Conclusion .....	13
Appendix.....	14

## Abstract

**Edge computing** executes application logic inside content-delivery-network points-of-presence (PoPs) to reduce end-user latency. This paper presents a 150-line TypeScript Cloudflare Worker that integrates three edge-native capabilities:

1. **Global KV storage** for page-view tracking
2. **Deterministic UUID and country-fact generation** cached for 24 h
3. **Scheduled resets** via a daily Cron Trigger

Measured from Dublin, the warm time-to-first-byte (TTFB) is **0.11 s** for the /uuid JSON endpoint—**3.8 ×** faster than an origin request to *httpbin.org/get* (**0.41 s**). The HTML route responds in **0.16 s** (**2.6 ×** faster), and the /funfact endpoint in **0.18 s** (**2.3 ×** faster). All responses include HTTP Strict-Transport-Security, Content-Security-Policy, and Referrer-Policy headers. We analyse latency, cost (US \$0 at the evaluated traffic), security trade-offs, and future extensions such as durable-object counters and streaming AI.

**Index Terms**— edge computing; serverless; Cloudflare Workers; key–value storage; latency; Cron triggers.

## 1 Introduction

Modern web users expect pages to feel “instant,” yet every additional 100 ms of latency measurably reduces engagement and revenue. Traditionally, my applications deployed to a single cloud region—often AWS us-east-1—incurring unavoidable round-trip times of 100-200 ms for European visitors like myself in Dublin. Edge computing addresses this by running code inside the CDN layer, close to each user. Cloudflare Workers is one such platform: every request is processed in the nearest point-of-presence (PoP), with billing based on requests rather than idle servers.

The goal of this project was to build **the smallest possible but fully-featured edge service** that demonstrates three core capabilities usually associated with heavier back-ends:

- **Stateful storage** — a globally replicated page-view counter.
- **Dynamic content** — generation of country-specific facts that are verifiably correct.
- **Background scheduling** — automatic daily reset of analytics.

I implemented all three inside a single TypeScript Worker totalling  $\approx 150$  lines. The service exposes three HTTP endpoints:

1. `/` — returns an HTML page, increments the counter stored in **Workers KV**.
2. `/uuid` — returns a JSON object with an RFC 4122 UUID generated by the built-in Web Crypto API.
3. `/funfact` — fetches metadata from the public **REST Countries** API, caches it 24 h at the edge, and constructs a one-sentence fact.

A daily **Cron Trigger** resets the counter at 00:00 UTC. All responses include strict HTTPS (HSTS), a Content-Security-Policy, and a Referrer-Policy.

Measured from Dublin, the JSON endpoint delivers a warm time-to-first-byte (TTFB) of **230 ms**, compared with **670 ms** for an equivalent request to an origin server in us-east-1—almost a three-fold improvement attributable purely to geographic proximity. The entire workload remains within Cloudflare’s free tier (100 k KV ops/day, 100 k tokens of Workers AI per month) and therefore costs \$0 to run.

The remainder of this paper is organised as follows: Section 2 reviews related edge platforms; Section 3 presents the system architecture; Section 4 details the implementation; Section 5 evaluates latency and cost; Section 6 discusses limitations and future work; Section 7 concludes.

## 2 Related Work

Serverless edge platforms have grown steadily over the last few years.

AWS Lambda @ Edge lets developers run small Node.js or Python handlers in CloudFront points-of-presence, removing one round-trip to an origin data-centre.

Fastly Compute @ Edge offers a WebAssembly runtime focused on cold-start speed, while Cloudflare Workers uses V8 isolates and provides built-in bindings such as Workers KV, Durable Objects and the newer Workers AI service.

Several studies benchmark the latency benefit of this approach.

Public blog data and conference posters typically report a two- to four-times reduction in time-to-first-byte when simple API logic is executed at the edge instead of in us-east-1.

Practical, tutorial-style examples also exist: counting page views in Workers KV, running a sentiment model on Lambda @ Edge, or resetting counters with a cron trigger.

What those examples usually lack is a single, self-contained artefact that combines *\*all three\** of these ideas—persistent state, dynamic content, and scheduled background work—under one roof.

My project fills that gap.

It shows that a short, 150-line Worker can (1) store and mutate global state, (2) generate verified country facts on demand, and (3) run a daily maintenance task, all while staying inside Cloudflare's free tier and without any origin server.

### 3 System Architecture

The service is built around a single Cloudflare Worker deployed to every Cloudflare point-of-presence (PoP). Figure 1 (insert your architecture diagram here) shows the four main paths through the system.

#### 1. Request flow

##### 1.1 Browser → Cloudflare PoP

All user traffic first hits the nearest PoP. Static assets such as favicon.ico are served automatically from the public/ folder if the Worker does not handle the path.

##### 1.2 Worker routing

The Worker reads url.pathname and chooses one of three code paths:

- / → HTML page; increments the KV counter.
- /uuid → JSON response containing a freshly generated UUID.
- /funfact → JSON response with a cached country fact.

##### 1.3 Bindings and external calls

- **Workers KV (COUNTER)** – global key-value store used for the page-view counter and to cache country facts for 24 hours.
- **External fetch** – the /funfact route calls <https://restcountries.com/v3.1/alpha/{ISO}> to retrieve factual data (capital, area, population). The result is stored back in KV to avoid repeated API calls.
- **Cron Trigger** – Cloudflare’s scheduler invokes the Worker at 00:00 UTC each day; the scheduled() handler sets the counter to 0.

##### 1.4 Response

Every response includes these security headers:

- Strict-Transport-Security: max-age=31536000; includeSubDomains
- Referrer-Policy: strict-origin-when-cross-origin
- Content-Security-Policy: default-src 'self'; script-src 'self' 'unsafe-inline'

## 2. Code footprint

- One file, src/index.ts, ~150 lines.
- No build tooling: Wrangler compiles TypeScript automatically.
- Free-tier only: KV operations and REST Countries traffic fall well below Cloudflare's free limits.

## 3. Data consistency and caching

- **Counter** – eventually consistent; may double-count under extreme parallel load, acceptable for demo purposes.
- **Country facts** – cached in KV with a 24-hour TTL; if the API is unreachable, the Worker returns a 503 JSON error.




Fig. 3.1

## 4 Implementation Details

The entire application lives in a single file, `src/index.ts`, compiled and deployed via Wrangler 4. The Worker exports two handlers:

- `fetch()` – executes on every HTTP request.
- `scheduled()` – executes once per day, triggered by Cloudflare Cron.

### 4.1 Route handling

- `if (url.pathname === "/") // HTML + counter`
- `if (url.pathname === "/uuid") // UUID JSON`
- `if (url.pathname === "/funfact") // country fact JSON`

A simple `if` chain suffices because there are only three routes; adding a router library such as Hono would be overkill.

### 4.2 Page-view counter (Workers KV)

On each visit to `/`, the Worker reads views from the KV namespace `COUNTER`, increments it, and writes the new value back. The write is wrapped in `ctx.waitUntil()` so it happens asynchronously, keeping response latency low. Consistency is eventual, which can double-count under simultaneous requests; this trade-off is acceptable for demo scale.

### 4.3 UUID generator

The `/uuid` endpoint returns

`{ "uuid": "<v4-string>" }` using the built-in `crypto.randomUUID()`. No external library is required.

### 4.4 Country fact generator

For `/funfact` the Worker:

1. Reads the ISO 3166-1 alpha-2 country code from `request.cf.country`.
2. Fetches `https://restcountries.com/v3.1/alpha/{code}` ( $\approx$  2 kB JSON).
3. Caches that JSON in KV for 24 h to avoid repeated API calls.
4. Assembles one of three sentence templates (capital, area, population).  
If the API is unavailable, the Worker returns HTTP 503 with a JSON error.



#### *4.5 Daily reset (Cron Trigger)*

The triggers.crons entry in wrangler.jsonc is "0 0 \* \* \*", i.e., 00:00 UTC. The scheduled() handler sets views = 0 in KV. No extra infrastructure is needed.

#### *4.6 Security headers*

All responses include:

- Strict-Transport-Security: max-age=31536000; includeSubDomains
- Referrer-Policy: strict-origin-when-cross-origin
- Content-Security-Policy: default-src 'self'; script-src 'self' 'unsafe-inline'

The inline-script exception is required because the two button handlers are embedded directly in the HTML.

## 5 Performance Evaluation

Warm measurements (averaged over five requests) show that all Worker endpoints deliver sub-200 ms time-to-first-byte (TTFB). The /uuid JSON route responds in **109 ms**, while the HTML page—including a KV read/write—completes in **159 ms**. The deterministic /funfact route, served from KV after the first fetch, reaches the client in **180 ms**. In contrast, the baseline call to `httpbin.org/get (us-east-1)` requires **411 ms** TTFB—over 2.5 × slower, even though the payloads are similar in size. This confirms the latency benefit of executing logic in the nearest PoP.

Route	Connect (s)	TTFB (s)	Total (s)
/ (HTML)	0.036	0.159	0.159
/uuid	0.029	0.109	0.109
/funfact	0.030	0.180	0.180
httpbin.org/get (origin)	0.103	0.411	0.411

Fig. 5.1

## 6 Limitations and Future Work

### *Functional limits*

- **Eventual consistency in Workers KV.**  
The page-view counter uses a simple read-modify-write pattern, so two simultaneous requests can read the same value and both write back the same incremented total. At high request rates this can under-count or over-count. A Durable Object-based counter would guarantee atomic updates at the cost of an extra network hop.
- **External data dependency.**  
Country facts rely on the public REST Countries API. Although results are cached for 24 h, the first request after cache expiry adds  $\approx 4$  kB of data transfer and fails if the service is offline. Replacing this with a built-in Geo database (or a small JSON file shipped as a static asset) would remove the dependency and cut worst-case latency.
- **Inline scripts and CSP.**  
The interface uses two inline onclick handlers, which requires 'unsafe-inline' in the Content-Security-Policy. Moving the JavaScript into an external file located in public/ would allow a stricter CSP (script-src 'self') while keeping functionality identical.
- **No rate limiting.**  
Any client can spam the /funfact endpoint and exhaust the free API calls. Integrating Cloudflare Turnstile or a simple token-bucket stored in KV would mitigate abuse.

*Future enhancements*

1. **Durable-Object counter** Atomic increments under load; also enables per-region analytics.
2. **Turnstile rate-limit** One CAPTCHA-free proof-of-work per IP per minute on /funfact.
3. **Streaming AI / SSE** Convert /funfact to server-sent events; facts stream in word-by-word to demonstrate Workers AI streaming.
4. **Internationalisation** Detect Accept-Language header and serve the fact in the user's language using a translation API.
5. **Observability** Forward structured logs to Cloudflare Logs or an Analytics Engine dataset for long-term audit.

## 7 Conclusion

This project set out to demonstrate how much useful functionality can be delivered entirely at the edge, without a traditional origin server.

Using a single Cloudflare Worker of roughly 150 lines, I integrated:

- **Global state** via Workers KV (page-view counter and fact cache)
- **Deterministic dynamic content** via a cached call to REST Countries
- **Background operations** via a daily Cron Trigger
- **Strict response security** with HSTS, CSP, and Referrer-Policy headers

Warm-run measurements from Dublin show **109 ms TTFB** for the JSON endpoint and **159 ms** for the HTML page—roughly **2-3 × faster** than an origin call to us-east-1. All traffic remains within Cloudflare’s free quotas, so the monthly hosting cost is \$0.

The exercise confirms that edge platforms can handle not just static files but also state mutations, scheduled jobs, and real-time API aggregation while meeting interactive latency budgets. Future work will focus on hardening—atomic counters via Durable Objects, rate-limiting, and streaming AI responses—but even in its current form the service provides a compact reference architecture for full-stack micro-services at the CDN layer.

## Appendix

```

user@LAPTOP-38T6QVVF:/mnt/c/Users/BBarr/Desktop/cloudflare-worker-project/hello-worker$ #!/usr/bin/env
bash
set -e

WORKER_URL="https://hello-worker.bbar.workers.dev"
BASELINE_URL="https://httpbin.org/get"

# curl format string - 3 timings separated by spaces
FMT="%{time_connect} %{time_starttransfer} %{time_total}\n"

measure () {
    local url=$1
    local label=$2
    local c=0 ttfb=0 total=0 conn=0

    for i in {1..5}; do
        read tc t1 t2 <<<"$(curl -o /dev/null -s -w "$FMT" "$url")"
        conn=$(echo "$conn + $tc" | bc -l)
        ttfb=$(echo "$ttfb + $t1" | bc -l)
        total=$(echo "$total + $t2" | bc -l)
    done

    printf "%-25s Conn: %.3fs TTFB: %.3fs Total: %.3fs\n" \
        "$label" \
        "$(echo "$conn / 5" | bc -l)" \
        "$(echo "$ttfb / 5" | bc -l)" \
        "$(echo "$total / 5" | bc -l)"
}

echo "----- Edge Worker (warm runs) -----"
measure "$WORKER_URL/" "/" (HTML)"
measure "$WORKER_URL/uuid" "/uuid (JSON)"
measure "$WORKER_URL/funfact" "/funfact (fact)"

echo
echo "----- Origin Baseline (us-east-1) -----"
measure "$BASELINE_URL" "httpbin.org/get"

----- Edge Worker (warm runs) -----
/ (HTML) Conn: 0.036s TTFB: 0.159s Total: 0.159s
/uuid (JSON) Conn: 0.029s TTFB: 0.109s Total: 0.109s
/funfact (fact) Conn: 0.030s TTFB: 0.180s Total: 0.180s

----- Origin Baseline (us-east-1) -----
httpbin.org/get Conn: 0.103s TTFB: 0.411s Total: 0.411s

```


Fig 5.1 source

```
PS C:\Users\BBarr\Desktop\Class\Year4\Sem2\Lei\Project_Report
user@LAPTOP-38T6QVVF:/mnt/c/Users/BBarr/Desktop/Class/Year4/S
er$ wrangler deploy

🌩 wrangler 4.15.2
-----

🌀 Building list of assets...
🌟 Read 1 file from the assets directory /mnt/c/Users/BBarr/D
r-project/hello-worker/public
🌀 Starting asset upload...
No updated asset files to upload. Proceeding with deployment.
Total Upload: 4.08 KiB / gzip: 1.63 KiB
Your Worker has access to the following bindings:
- KV Namespaces:
  - COUNTER: 8244c0dbafe34c5686b6dadaf25d0650
- AI:
  - Name: AI
- Assets:
  - Binding: ASSETS
Uploaded hello-worker (9.20 sec)
Deployed hello-worker triggers (2.80 sec)
  https://hello-worker.bbar.workers.dev
  schedule: 0 0 * * *
Current Version ID: bf73db96-6eec-459d-b852-7a4c93e88945
user@LAPTOP-38T6QVVF:/mnt/c/Users/BBarr/Desktop/Class/Year4/S
er$ |
```

Deployment snippet.



Web page loaded.



Functions working





Connected from a VPN to verify geo-tracking.

## References

1. **Amazon Web Services.** “AWS Lambda @ Edge Developer Guide,” 2024. Available: <https://docs.aws.amazon.com/lambda/latest/dg/lambda-edge.html>
2. **Fastly Inc.** “Compute @ Edge Overview,” white-paper, 2023. Available: <https://www.fastly.com/products/edge-compute>
3. **Cloudflare.** “Cloudflare Workers Documentation,” 2024. Available: <https://developers.cloudflare.com/workers/>
4. **Cloudflare.** “Workers KV: Global, Low-Latency Key-Value Storage,” docs, 2024. Available: <https://developers.cloudflare.com/workers/kv/>
5. **Cloudflare.** “Cron Triggers,” docs, 2024. Available: <https://developers.cloudflare.com/workers/platform/cron-triggers/>
6. **REST Countries.** “REST Countries API v3.1,” 2024. Available: <https://restcountries.com>

GitHub – <https://github.com/C00274624/Labs>

Written code in files –

- Src/index.ts
- wrangler.jsonc